

000
001
002
003
004
005
006
007
008
009
010
011
012
013
014
015
016
017
018
019
020
021
022
023
024
025
026
027
028
029
030
031
032
033
034
035
036
037
038
039
040
041
042
043
044
045
046
047
048
049
050
051
052
053

Bayesian Modeling

Contents

1 Overview	1
1.1 Goal	1
1.2 Target Audience	1
1.3 Prerequisites	2
1.4 Format	2
2 Topics	2
2.1 Introduction to Bayes	2
2.2 Methods	3
2.3 More complicated models	3
2.4 Even more complicated models	3
2.5 Bayesian workflow	3
2.6 Student project presentations	4
A Resources which may be appropriate for mini-reading group or student presentations	4
B Leads on concrete exercises / projects	4
B.1 Batting average dataset	4
C Need to do	5

1 Overview

1.1 Goal

The goal of this workshop is to introduce students to the concepts and practice of Bayesian modeling.

1.2 Target Audience

We expect that the typical student will be a graduate student, faculty member, staff member, or researcher in a quantitative field (such as computer science, statistics, engineering, or biology), who would like to learn more about Bayesian modeling.

1.3 Prerequisites

Prerequisites include calculus, some linear algebra, and some familiarity with introductory probability (e.g., we will assume prior familiarity with concepts such as expectation, conditional probability, and commonly used distributions, such as Gaussian and Poisson.) We will work in Python when working collaboratively as a class, but you can do offline analyses in any language of your choosing.

1.4 Format

We will present lectures via slides. But I would like there to be student contribution, active learning (in a non-artificial sense), and some student autonomy and self-direction.

To this end, I am leaning leaving the last chunk of the workshop open for students to present a mini-project they did on Bayesian data analysis. Probably this would be chosen after covering the material in Section 2.3.

That said, along the way, we could also potentially try to encourage student involvement via discussions and active-learning components embedded within the material, e.g.

- Mini reading group discussions – They might not have time to read a whole paper, but I could assign sections of a paper, or sections of a text, that we discuss at the very beginning.
- Student “lightning chat” (10 minute) presentations? Could do one per student per workshop. These could consist of students choosing any of the following:
 - Presentation of Python implementations of models from [1] , [2], or the workshop.
 - Presentation of an exercise from [2] or [1].
 - Presentation of a reading section, blog, etc. of interest.
 - Presentation of a mathematical derivation of something relevant to the course.
 - Presentation of how a concept relates to something from their research area.
 - Presentation of mid-progress on their student project.
 - A list of questions that we can discuss.
- Real-time python applications lab – Google Collab exercises ? Python (rather than R) implementations of [1] and [2] ? Progress on their project? Etc.

2 Topics

Below are topics we plan to cover in the course:

2.1 Introduction to Bayes

We present everything in here using conjugate models with closed-form posteriors. The models are useful in and of themselves, as well as to build intuition for more complicated models.

Primary references here are [1] and [2].

- **Why Bayes?** – See Section 1.3 of [1]. [3] has some nice plots motivating why use Bayesian linear regression over standard linear regression. [4] has some nice plots illustrating the Bayesian approach and how it mitigates overfitting. I can provide a nice example with biometric profiling of human typing dynamics. [5] has a nice simple example of obtaining non-standard functionals from the posterior that can be of interest. [6] presents the case for Bayesian deep learning.
- **Belief functions, Bayes rule** – Sections 2.1, 2.2 of [1]. [4] briefly overviews of the Bayesian framework. *Why most published research findings are false* [7] provides nice motivation. Could perhaps cover exchangeability here.
- **Binomial, Poisson, normal, multivariate normal models** – Sections 3.1, 3.2, 5, and 5 of [1]. Introduce the exponential family formalism [8] for much greater breadth.

- **Bayesian linear regression** – Section 9 of [1]. I have notes on this. There are some nice slides here which also illustrate the use of kernels.¹ Introduce model selection here (Section 9.3 of [1]).

We will want to find a way to get students to group up, probably based on domain expertise/interests, so that they can eventually work together on a project.

2.2 Methods

We introduce these methods, which can be used for models without closed-form posteriors. We practice applying them in the next section.

- **MCMC** - [Karin](#) will present.
- **Variational inference** [9].

2.3 More complicated models

Here are some models which are still fairly standard, but lack conjugate priors, and so inference typically requires VI or MCMC. [Karin: Where here, or elsewhere, would you like to illustrate applications of MCMC?](#)

- **Hierarchical models** Hierarchical normal model (e.g. Gelman's 5 schools example), hierarchical linear regression (Chapter 13 of [2]).
- **Regression models for binary and multi-class data** Includes logistic regression, probit regression, binomial, multinomial, etc. Use this to cover additional inference techniques: auxiliary variable trick and Laplace variational inference. See also pp. 390 of [1] for a useful warm-starting strategy. Could generalize to Bayesian GLMs. Could also cover or mention hierarchical extensions (i.e. Bayesian GLMM's).
- **Mixture models** I will give CAVI for Gaussian mixture models.
- **Time series models** Probably just hidden markov models, although would be nice to also introduce state space models. Could mention embedding of GLM's or GLMM's within them. May give some overview to probabilistic graphical models here.

2.4 Even more complicated models

Here we discuss models which have additional complexity – they could involve neural networks, non-parametrics, larger scale, etc. Often these involve some additional inferential machinery – *stochastic* variational inference, reparametrization trick, etc.

- **Why Bayesian Deep Learning?** Bayes and neural networks. 20-30 min w/ guest presenter, Kyle Heuton, Ph.D. student, computer science.
- **Custom models** Could present black-box variational inference or automatic differentiation variational inference here. I have some notes that I could convert to slides. Would take some work though. [Anna, perhaps you'd be interested, since you've been working with this?](#)
- **Sampler platter of other possible topics** VAE, Gaussian processes, composing time series models with neural networks, indian buffet processes, dirichlet process mixture models, etc.

2.5 Bayesian workflow

Lots of nice resources for Bayesian workflow. For example: [10] or [11]. Section 6 of [2] covers model checking.

¹Nice Bayesian linear regression slides: https://www.cs.toronto.edu/~rgrosse/courses/csc411_f18/slides/lec19-slides.pdf

We could cover this the day before the student projects – which should help them as they wrap up their projects – and leave lots of time/space for workshopping.

Some points to make re: model checking

- *Samples from the posterior predictive should capture important properties of the observed dataset.* For a violation of this, see the normal model for Newcomb’s speed of light measurements. (Compare Figures 6.2 and 3.1 of [2].)

2.6 Student project presentations

For student projects, we could have students present results of a Bayesian data analysis mini-project. Perhaps they could highlight some aspect of the Bayesian workflow along the way.

We would probably want to have a “workshopping” section the day before they do their presentations.

References

- [1] Peter D Hoff. *A first course in Bayesian statistical methods*, volume 580. Springer, 2009.
- [2] Andrew Gelman, John B Carlin, Hal S Stern, David B Dunson, Aki Vehtari, and Donald B Rubin. *Bayesian data analysis*. CRC press, 2013.
- [3] Christopher M Bishop. *Pattern recognition and machine learning*. springer, 2006.
- [4] Zoubin Ghahramani. Bayesian non-parametrics and the probabilistic approach to modelling. *Philosophical Transactions of the Royal Society A: Mathematical, Physical and Engineering Sciences*, 371(1984):20110553, 2013.
- [5] Leonhard Held and Chris C Holmes. Bayesian auxiliary variable models for binary and multinomial regression. *Bayesian analysis*, 1(1):145–168, 2006.
- [6] Andrew Gordon Wilson. The case for bayesian deep learning. *arXiv preprint arXiv:2001.10995*, 2020.
- [7] John PA Ioannidis. Why most published research findings are false. *PLoS medicine*, 2(8):e124, 2005.
- [8] Michael Wojnowicz. *The exponential family*. Available (with permission).
- [9] Michael Wojnowicz. *Foundations of variational inference*. Available (with permission) at https://github.com/mikewojnowicz/vi_foundations.
- [10] Andrew Gelman, Aki Vehtari, Daniel Simpson, Charles C Margossian, Bob Carpenter, Yuling Yao, Lauren Kennedy, Jonah Gabry, Paul-Christian Bürkner, and Martin Modrák. Bayesian workflow. *arXiv preprint arXiv:2011.01808*, 2020.
- [11] Jonah Gabry, Daniel Simpson, Aki Vehtari, Michael Betancourt, and Andrew Gelman. Visualization in bayesian workflow. *Journal of the Royal Statistical Society: Series A (Statistics in Society)*, 182(2):389–402, 2019.

A Resources which may be appropriate for mini-reading group or student presentations

- Intro: *Why most published research findings are false*. [7]
- Model checking: Bayesian workflow.
- Textbook sections TBD.

B Leads on concrete exercises / projects

B.1 Batting average dataset

The hierarchical normal model for (arcsine-transformed) batting average data on pp. 163 of [2] has some serious deficiencies, as exposed in Table 6.1 in the section on model checking.

216 Can you construct (and learn) a better model which makes predictions closer to the true final batting
217 average?

218
219 Examples:

- 220 • Add an extra layer to the hierarchy, so that player p 's 1970 batting average inherits from
221 player p 's overall batting average which in turn inherits from a population batting average.
222 (Of course, I am speaking of the arcsine-transformed batting averages, so that we can use
223 a hierarchical normal model.)
- 224 • Add an autoregressive component, because, as mentioned by Gelman, player batting aver-
225 ages *DO* change over time.

226
227 The text also does a poor job of checking the modeling assumption violations that were of concern.
228 Can you do a better job of checking them, and if necessary, address them?

229
230 Examples:

- 231 • If batting averages are indeed heavy tailed or skewed, move from a normal distribution to
232 something else. For example, could try a t-distribution with Laplace inference to handle
233 the non-conjugacy.
- 234 • If the variance is indeed too high for a binomial model, try something that can handle the
235 overdispersion.

236 237 **C Need to do**

- 238
239 • [Karin](#) Bring pymc3, stan, etc. into this – will make things a lot more useful for the audience
240 than requiring that they code things up from scratch!

241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269